

cem: Coarsened Exact Matching in Stata

Matthew Blackwell¹ Stefano Iacus² Gary King³ Giuseppe Porro⁴

September 21, 2014

¹Institute for Quantitative Social Science, 1737 Cambridge Street, Harvard University, Cambridge MA 02138; <http://www.mattblackwell.org>, mblackwell@gov.harvard.edu).

²Department of Economics, Business and Statistics, University of Milan, Via Conservatorio 7, I-20124 Milan, Italy; stefano.iacus@unimi.it.

³Institute for Quantitative Social Science, 1737 Cambridge Street, Harvard University, Cambridge MA 02138; <http://GKing.Harvard.Edu>, King@Harvard.Edu, (617) 495-2027).

⁴Department of Economics and Statistics, University of Trieste, P.le Europa 1, I-34127 Trieste, Italy; giuseppe.porro@econ.units.it.

Abstract

This paper introduces a Stata implementation of Coarsened Exact Matching (CEM), a new method for improving the estimation of causal effects by reducing imbalance in covariates between treated and control groups. CEM is faster, easier to use and understand, requires fewer assumptions, more easily automated, and possesses more attractive statistical properties for many applications than existing matching methods. In CEM, users temporarily coarsen their data, exact match on these coarsened data, then run their analysis on the uncoarsened, matched data. CEM bounds the degree of model dependence and causal effect estimation error by ex ante user choice, is monotonic imbalance bounding (so that reducing the maximum imbalance on one variable has no effect on others), does not require a separate procedure to restrict data to common support, meets the congruence principle, is approximately invariant to measurement error, balances all nonlinearities and interactions in-sample (i.e., not merely in expectation), and works with multiply imputed data sets. Other matching methods inherit many of CEM's properties when applied to further match data preprocessed by CEM. The library `cem` implements the CEM algorithm in Stata.

1 Introduction

This program is designed to improve the estimation of causal effects via a powerful method of matching that is widely applicable in observational data and easy to understand and use (if you understand how to draw a histogram, you will understand this method). The program implements the Coarsened Exact Matching (CEM) algorithm described in (Iacus, King and Porro, 2012). CEM is a monotonic imbalance reducing matching method — which means that the balance between the treated and control groups is chosen by ex ante user choice rather than discovered through the usual laborious process of checking after the fact, tweaking the method, and repeatedly reestimating. CEM also assures that adjusting the imbalance on one variable has no effect on the maximum imbalance of any other. CEM strictly bounds through ex ante user choice both the degree of model dependence and the average treatment effect estimation error, eliminates the need for a separate procedure to restrict data to common empirical support, meets the congruence principle, is robust to measurement error, works well with multiple imputation methods for missing data, can be completely automated, and is extremely fast computationally even with very large data sets. After preprocessing data with CEM, the analyst may then use a simple difference in means or whatever statistical model they would have applied without matching. CEM can also be used to improve other methods of matching by applying those methods to CEM-matched data (they formally inherit CEM’s properties if applied within CEM strata). CEM also works well for determining blocks in randomized experiments, and evaluating extreme counterfactuals.

2 Background

2.1 Notation

Consider a sample of n units randomly drawn from a population of N units, where $n \leq N$. For unit i , denote T_i as an indicator variable with value $T_i = 1$ if unit i receives the treatment (and so is a member of the “treated” group) and $T_i = 0$ if not (and is therefore a member of the “control” group). The outcome variable is denoted Y , where $Y_i(0)$ is the potential outcome for observation i if the unit does not receive treatment and $Y_i(1)$ is the potential outcome if the (same) unit receives treatment. For each observed unit, the observed outcome is $Y_i = T_i Y_i(1) + (1 - T_i) Y_i(0)$ and so $Y_i(0)$ is unobserved if i receives treatment and $Y_i(1)$ is unobserved if i does not receive treatment.

To compensate for the observational data problem where the treated and control groups are not necessarily identical before treatment (and, lacking random assignment, not the same on average), matching estimators attempt to control for pre-treatment covariates. For this purpose, we denote $X = (X_1, X_2, \dots, X_k)$ as a k -dimensional data set, where each X_j is a column vector of observed values of pre-treatment variable j for the n sample observations (possibly drawn from a population, of size N). That is, $X = [X_{ij}, i = 1, \dots, n, j = 1, \dots, k]$.

2.2 Quantities of Interest

As usual, the treatment effect for unit i , $TE_i = Y_i(1) - Y_i(0)$, is unobserved. All relevant causal quantities of interest are functions of TE_i , for different groups of units, and so must

be estimated. We focus on the sample average treatment effect on the treated (SATT):

$$\text{SATT} = \frac{1}{n_T} \sum_{i \in T} \text{TE}_i \quad (1)$$

where $n_T = \sum_{i=1}^n T_i$ and $T = \{1 \leq i \leq n : T_i = 1\}$. Matching algorithms sometimes also change the quantity being estimated to one that can be estimated without much model dependence by selecting control and/or treated units.

We assume that treatment assignment is ignorable conditional on X . This assumption is often stated as “no unmeasured confounders” or “no omitted variables.” Formally, this means that the treatment assignment is independent of the potential outcomes,

$$P(T|X, Y(0), Y(1)) = P(T|X). \quad (2)$$

2.3 Existing matching methods and practice

Matching is a nonparametric method of controlling for some or all of the confounding influence of pretreatment control variables in observational data. The key goal of matching is to prune observations from the data so that the remaining data have better *balance* between the treated and control groups, meaning that the empirical distributions of the covariates (X) in the groups are more similar. Exactly balanced data means that controlling further for X is unnecessary (since it is unrelated to the treatment variable), and so a simple difference in means on the matched data can estimate the causal effect; approximately balanced data requires controlling for X with a model (such as the same model that would have been used without matching), but the only inferences necessary are only those relatively close to the data, leading to less model dependence and reduced statistical bias than without matching.

The most common matching methods involve finding, for each treated unit, at least one control unit that is “similar” on the covariates. The distinction between methods is how to define this similarity. For example, *exact matching* simply matches a treated unit to all of the control units with the same covariate values. Unfortunately, due to the richness of covariates in many examples, this method often produces very few matches. A whole host of *approximate matching* methods specify a metric to find control units that are close to the treated unit. This metric is often the Mahalanobis distance or the propensity score (which is simply the probability of being treated, conditional on the covariates). Many of these related methods are implemented in Stata (Becker and Ichino, 2002; Abadie et al., 2004; Leuven and Sianesi, 2004; Abadie, Diamond and Hainmueller, 2009, forthcoming). A problem with this type of solution is that it requires the user to set the size of the matching solution ex ante, then check for balance ex post. Thus analysts must check for balance after the algorithm is finished, then respecify a matching model and recheck balance, etc. This process repeats until the user obtains an acceptable amount of balance.

As matching is simply a data preprocessing technique, analysts must still apply statistical estimators to the data after matching. When one-to-one exact matching is used, a simple difference in means between Y in the treated and control group provides an estimator of the causal effect. When the match is not exact, a parametric model must be used to control for the differences in the covariates across treated and control groups. This may be a linear regression, a maximum likelihood estimator or some other estimator. Applying a matching method to the data before analysis can reduce the degree of model dependence (Ho et al., 2007).

One wrinkle in the analysis of matched data occurs when there are not equal numbers of treated and control units within strata. In this situation, estimators require weighting observations according to the size of their strata (Iacus, King and Porro, 2012).

3 Coarsened Exact Matching

3.1 The Algorithm

The central motivation for CEM is that while exact matching provides perfect balance, it typically produces few matches due to curse-of-dimensionality issues. For instance, adding one continuous variable to a dataset effectively kills exact matching since two observations are unlikely to have identical values on a continuous measure. The idea of CEM is to temporarily coarsen each variable into substantively meaningful groups, exact match on these coarsened data and then only retain the original (uncoarsened) values of the matched data. As coarsening is a process at the heart of measurement, many analysts know how to coarsen a variable into groups that preserve information. For instance, education may be measured in years, but many would be comfortable grouping observations into categories of high school, some college, college graduates, etc. This method works by exact matching on distilled information in the covariates as chosen by the user.

The algorithm works as follows:

1. Begin with the covariates X and make a copy, which we denote X^* .
2. Coarsen X^* according to user-defined cutpoints, or CEM's automatic binning algorithm.
3. Create one stratum per unique observation of X^* and place each observation in a stratum.
4. Assign these strata to the original data, X and drop any observation whose stratum does not contain at least one treated and one control unit.

Once completed, these strata are the foundations for calculating the treatment effect. The inherent trade-off of matching is reflected in CEM too: larger bins (more coarsening) used to make X^* will result in fewer strata. Fewer strata will result in more diverse observations within the same strata and, thus, higher imbalance.

It is important to note that CEM prunes both treated and control units. This process changes the quantity of interest under study to the treatment effect in the post-matching subsample. This change is reasonable so long as the decision is transparent (see e.g. Crump et al. (2006)).

3.2 The Benefits

Iacus, King and Porro (2012) derive many of the properties of the CEM algorithm and we review some of them here. The key property of CEM is that it is in a class of matching methods called Monotonic Imbalance Bounding (MIB). MIB methods bound the maximum imbalance in some feature of the empirical distributions through an *ex ante* choice by the user. In CEM, this *ex ante* choice is the coarsening. As the coarsening on any variable becomes finer (the bins become more narrow), the bound on the maximum imbalance on the moments of that variable becomes tighter. This is also true for the bound on differences in the empirical quantiles. Furthermore, this choice also bounds the maximum imbalance on the full multivariate histogram of treated and control units, which includes

all interactions and non-linearities. By choosing the coarsening *ex ante*, users can control the amount of imbalance in the matching solution. Iacus, King and Porro (2012) also show that CEM bounds both the error in estimating the average treatment effect and the amount of model dependence.

Aside from bounding the imbalance between the treated and control groups, CEM has a number of other beneficial properties. First, CEM meets the *congruence principle*, which states that the data space and analysis space should be the same. Methods that fail to meet this principle often produce strange or counter-intuitive results. Methods that meet the principle allow analysts to leverage their substantive knowledge of the data in order to find better matches. Second, CEM automatically restricts the matched data to areas of common empirical support. This is necessary to remove the possibility of difficult-to-justify extrapolations of the causal effect that end up being heavily model dependent (King and Zeng, 2006). Finally, CEM is computationally very efficient even for large data sets.

4 An Extended Example

We show here how to use CEM¹ through a simple running example: the National Supported Work (NSW) Demonstration data, also known as the Lalonde data set (Lalonde, 1986). This program provided training to selected individuals for 12-18 months and help finding a job in the hopes of increasing their earnings. The treatment variable, `treated`, is 1 for participants (the treatment group) and 0 for nonparticipants (the control group). The key outcome variable is earnings in 1978 (`re78`). The statistical goal is to estimate a specific version of a causal effect: the sample average treatment effect on the treated (the “SATT”).

Since participation in the program was not assigned strictly at random, we must control for a set of pretreatment variables by the CEM algorithm. These pre-treatment variables include age (`age`), years of education (`education`), marital status (`married`), lack of a high school diploma (`nodegree`), race (`black`, `hispanic`), indicator variables for unemployment in 1974 (`u74`) and 1975 (`u75`), and real earnings in 1974 (`re74`) and 1975 (`re75`). Some of these are dichotomous (`married`, `nodegree`, `black`, `hispanic`, `u74`, `u75`), some are categorical (`age` and `education`), and the earnings variables are continuous and highly skewed with point masses at zero. You can load this data into Stata using the command

```
use http://www.mattblackwell.org/files/stata/data/lalonde.dta, clear
```

Matching is not a method of estimation; it is a way to preprocess a data set so that estimation of SATT based on the matched data set will be less “model-dependent” (i.e., less a function of apparently small and indefensible modeling decisions) than when based on the original full data set. Matching involves pruning observations that have no close matches on pre-treatment covariates in both the treated and control groups. The result is typically less model-dependence, lower bias, and (by removing heterogeneity) increased efficiency (King and Zeng, 2006; Ho et al., 2007; Iacus, King and Porro, 2012).

4.1 Basic Evaluation and Analysis of Unmatched Data

We begin the simple difference in means as a naive estimate of SATT; this estimator is useful only when the in-sample distribution of pre-treatment covariates happens to be the

¹In addition to the Stata version of CEM, there is an R version in the package `cem`. The example presented here is also used in that package as a vignette, and includes some obvious overlap in prose.

same in the treatment and control groups. First we compute the size of the treated and control groups:

```
. table treated
```

treated	Freq.
0	425
1	297

Thus, the data include 297 treated units and 425 control units. The (unadjusted and therefore likely biased) difference in means can be found by a simple linear regression of outcome on treatment,

```
. regress re78 treated
```

Source	SS	df	MS	Number of obs	=	722
Model	137332528	1	137332528	F(1, 720)	=	3.52
Residual	2.8053e+10	720	38962865.4	Prob > F	=	0.0609
Total	2.8191e+10	721	39099300.5	R-squared	=	0.0049
				Adj R-squared	=	0.0035
				Root MSE	=	6242

re78	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
treated	886.3038	472.0863	1.88	0.061	-40.52625 1813.134
_cons	5090.048	302.7826	16.81	0.000	4495.606 5684.491

Thus, our estimate of SATT is 886.3. Because the variable `treated` was not randomly assigned, the pre-treatment covariates differ between the treated and control groups. To see this, we focus on these pre-treatment covariates: `age`, `education`, `black`, `nodegree`, `re74`.

The overall imbalance is given by the \mathcal{L}_1 statistic, introduced in Iacus, King and Porro (2012) as a comprehensive measure of global imbalance. It is based on the L_1 difference between the multidimensional histogram of all pretreatment covariates in the treated group and that in the control group. First, we coarsen the covariates into bins. To use this measure, we require a list of bin sizes for the numerical variables. Our functions compute these automatically, or they can be set by the user.² Then, we cross-tabulate the discretized variables as $X_1 \times \dots \times X_k$ for the treated and control groups separately,

²Of course, as with drawing histograms, the choice of bins affects the final result. The crucial point is to choose one and keep it the same throughout to allow for fair comparisons. The particular choice is less crucial.

and record the k -dimensional relative frequencies for the treated $f_{\ell_1 \dots \ell_k}$ and control $g_{\ell_1 \dots \ell_k}$ units. Finally, our measure of imbalance is the absolute difference over all the cell values:

$$\mathcal{L}_1(f, g) = \frac{1}{2} \sum_{\ell_1 \dots \ell_k} |f_{\ell_1 \dots \ell_k} - g_{\ell_1 \dots \ell_k}| \quad (3)$$

Perfect global balance (up to coarsening) is indicated by $\mathcal{L}_1 = 0$, and larger values indicate larger imbalance between the groups, with a maximum of $\mathcal{L}_1 = 1$, which indicates complete separation. If we denote the relative frequencies of a matched dataset by f^m and g^m , then a good matching solution would produce a reduction in the \mathcal{L}_1 statistic; that is, we would hope to have $\mathcal{L}_1(f^m, g^m) \leq \mathcal{L}_1(f, g)$.

We compute \mathcal{L}_1 statistic, as well as several unidimensional measures of imbalance via our `imb` function. In our running example:

```
. imb age education black nodegree re74, treatment(treated)
```

```
Multivariate L1 distance: .50759358
```

```
Univariate imbalance:
```

	L1	mean	min	25%	50%	75%	max
age	.10119	.1792	0	1	0	-1	-6
education	.10047	.19224	1	0	1	1	2
black	.00135	.00135	0	0	0	0	0
nodegree	.08348	-.08348	0	-1	0	0	0
re74	.0522	-101.49	0	0	69.731	584.92	-2139

Only the overall \mathcal{L}_1 statistic measure includes imbalance with respect to the full joint distribution, including all interactions, of the covariates; in the case of our example, $\mathcal{L}_1 = .5076$. The \mathcal{L}_1 value is not valuable on its own, but rather as a point of comparison between matching solutions. The value .5076 is a baseline reference for the unmatched data. Once we have a matching solution, we will compare its \mathcal{L}_1 value to .5076 and gauge the increase in balance due to the matching solution from that difference. Thus, \mathcal{L}_1 works for imbalance as R^2 works for model fit: the absolute values mean less than comparisons between matching solutions. The unidimensional measures in the table are all computed for each variable separately.

The first column, labeled **L1**, reports the \mathcal{L}_1^j measure, which is \mathcal{L}_1 computed for the j -th variable separately (which of course does not include interactions). The second column in the table of unidimensional measures, labeled **mean**, reports the difference in means. The remaining columns in the table report the difference in the empirical quantiles of the distributions of the two groups for the 0th (min), 25th, 50th, 75th, and 100th (max) percentiles for each variable.

This particular table shows that variables **re74** is imbalanced in the raw data in many ways and variable **age** is balanced in means but not in the quantiles of the two distributions. This table also illustrates the point that balancing only the means between the treated and control groups does not necessarily guarantee balance in the rest of the distribution. Most important, of course, is the overall \mathcal{L}_1 measure, since even if the marginal distribution of every variable is perfectly balanced, the joint distribution can still be highly imbalanced.

4.2 Coarsened Exact Matching

We now apply the coarsened exact matching algorithm by calling the function `cem`. The CEM algorithm performs exact matching on coarsened data to determine matches and then passes on the uncoarsened data from observations that were matched to estimate the causal effect. Exact matching works by first sorting all the observations into strata, each of which has identical values for all the coarsened pre-treatment covariates, and then discarding all observations within any stratum that does not have at least one observation for each unique value of the treatment variable.

To run this algorithm, we must choose a type of coarsening for each covariate. We show how this is done via a fully automated procedures in next section. Then we show how to use explicit prior knowledge to choose the coarsening, which is normally preferable when feasible.

In CEM, the treatment variable may be *dichotomous* or *multichotomous*³. Alternatively, `cem` may be used for *randomized block experiments* without specifying a treatment variable; in this case the strata are simply returned without any pruning of observations.

4.2.1 Automated Coarsening

In our running example we have a dichotomous treatment variable. In the following code, we match on our chosen pre-treatment variables, but not `re78`, which is the outcome variable and so should never be included.

The output contains useful information about the match, including a (small) table about the number of observations in total, matched, and unmatched by treatment group, as well as the results of a call to the `imb` function for information about the quality of the matched data. Since `cem` bounds the imbalance ex ante, the most important information is the number of observations matched. But the results also give the imbalance in the matched data using the same measures as that in the original data described in Section 4.1. Thus,

```
. cem age education black nodegree re74, tr(treated)
```

Matching Summary:

Number of strata: 205

Number of matched strata: 67

	0	1
All	425	297
Matched	324	228
Unmatched	101	69

Multivariate L1 distance: .46113967

Univariate imbalance:

³ While CEM can match for multichotomous treatments, analysis with these matched samples is somewhat difficult. For instance, Iacus, King and Porro (2012) develop weights for two treatment groups and it is not obvious how to generalize these weights for more treatment groups. We suggest users run CEM on each pair of treatment levels, get the correct weights for each and calculate separate ATT.

	L1	mean	min	25%	50%	75%
age	.13641	-.17634	0	0	0	0
education	.00687	.00687	1	0	0	0
black	3.2e-16	-2.2e-16	0	0	0	0
nodegree	5.8e-16	4.4e-16	0	0	0	0
re74	.06787	34.438	0	0	492.23	39.425

	max
age	-1
education	0
black	0
nodegree	0
re74	96.881

We can see from these results the number of observations matched and thus retained, as well as those which were pruned because they were not comparable. By comparing the imbalance results to the original imbalance table given in the previous section, we can see that a good match can produce a substantial reduction in imbalance, not only in the means, but also in the marginal and joint distributions of the data.

The function `cem` also generates weights for use in the evaluation of imbalance measures and estimates of the causal effect (stored in `cem.weights`).

4.2.2 Coarsening by Explicit User Choice

The power and simplicity of CEM comes from choosing the coarsening yourself rather than using the automated algorithm as in the previous section. Choosing the coarsening enables you to set the maximum level of imbalance ex ante, which is a direct function of the coarsening you choose. By controlling the coarsening, you also put an explicit bound on the degree of model dependence and the SATT estimation error.

Fortunately, the coarsening is a fundamentally substantive act, almost synonymous with the measurement of the original variables. In other words, if you know something about the data you are analyzing, you almost surely have enough information to choose the coarsening. (And if you don't know something about the data, you might ask why you are analyzing it in the first place!)

In general, we want to set the coarsening for each variable so that substantively indistinguishable values are grouped and assigned the same numerical value. Groups may be of different sizes if appropriate. Recall that any coarsening during CEM is used only for matching; the original values of the variables are passed on to the analysis stage for all matched observations.

For numerical variables, we can use the cutpoints syntax in `cem`. Thus, for example, in the US educational system, the following discretization of years of education corresponds to different levels of school

Grade school	0–6
Middle school	7–8
High school	9–12
College	13–16
Graduate school	>16

Using these natural breaks in the data to create the coarsening is generally a good approach and certainly better than using fixed bin sizes (as in caliper matching) that disregard these meaningful breaks. In our data, no respondents fall in the last category,

```
. table education
```

education	Freq.
3	1
4	6
5	5
6	7
7	15
8	62
9	110
10	162
11	195
12	122
13	23
14	11
15	2
16	1

We can use the cutpoints above using parentheses after the `education` variable:

```
. cem age education (0 6.5 8.5 12.5 17.5) black nodegree re74, tr(treated)
```

Matching Summary:

```
Number of strata: 155
Number of matched strata: 53
```

	0	1
All	425	297
Matched	349	245
Unmatched	76	52

Multivariate L1 distance: .43604654

Univariate imbalance:

	L1	mean	min	25%	50%	75%	max
age	.05034	-.15556	0	0	0	1	-1
education	.0309	.00362	1	-1	0	0	2
black	8.2e-16	1.0e-15	0	0	0	0	0
nodegree	1.2e-15	1.9e-15	0	0	0	0	0

```
re74    .04975    2.5048        0        0    161.28    -17.37    1198.1
```

As we can see, this matching solution differs from that resulting from our automated approach in the previous section. In fact, it has actually increased the balance in matching solution while giving us a higher number of matched units.

4.2.3 Coarsening categorical variables

For categorical variables that do not have a natural ordering, some recoding might be necessary before inputting to CEM. For instance, if we have a variable that is

Strongly Agree	1
Agree	2
Neutral	3
Disagree	4
Strongly Disagree	5
No Opinion	6

there is a category (“No Opinion”) that does not fit on the ordinal scale of the variable. In our example dataset, we have such a variable, `q1`,

```
. table q1
```

q1	Freq.
strongly agree	121
agree	111
neutral	129
disagree	121
strongly disagree	118
no opinion	122

In order to coarsen this variable, first create a new coarsened variable using the `recode` command⁴:

```
. recode q1 (1 2 = 1 "agree") (3 6 = 2 "neutral") (4 5 = 3 "disagree"), gen(cem_q1)
(601 differences between q1 and cem_q1)

. table cem_q1
```

RECODE of	
q1	Freq.
1	
2	
3	

⁴For variables that are strictly string (non-numeric) variables, users will need to first use the `encode` command to convert the strings to numeric, then use `recode`.

agree	232
neutral	251
disagree	239

Here we have collapsed the opinions into the direction of opinion, also grouping “No Opinion” with “Neutral.” Once the coarsened variable is created, you can pass this variable to CEM with the (#0) cutpoints command after it to ensure that CEM does not coarsen further:

```
. cem age education black nodegree re74 cem_q1 (#0), tr(treated)
```

Matching Summary:

Number of strata: 315

Number of matched strata: 81

	0	1
All	425	297
Matched	260	190
Unmatched	165	107

Multivariate L1 distance: .5904067

Univariate imbalance:

	L1	mean	min	25%	50%	75%	max
age	.14574	-.1994	0	0	0	1	-1
education	.00263	.00263	1	0	0	0	0
black	3.6e-16	6.7e-16	0	0	0	0	0
nodegree	3.5e-16	6.7e-16	0	0	0	0	0
re74	.09854	70.061	0	0	375.1	-383.76	96.881
cem_q1	3.1e-16	3.1e-15	0	0	0	0	0

When calculating treatment effects after running CEM, be sure to use the original, uncoarsened variables for analysis. Coarsened variable should only be used to produce matches. After this, they can be discarded.

4.3 Restricting the matching solution to a k -to- k match

By default, CEM uses maximal information, resulting in strata that may include different numbers of treated and control units. To compensate for the differential strata sizes, `cem` also returns weights to be used in subsequent analyses. Although this is generally the best option, a user with enough data may opt for a k -to- k solution to avoid the slight inconvenience of needing to use weights.

The argument `k2k` accomplishes this by pruning observations from a `cem` solution within each stratum until the solution contains the same number of treated and control

units within all strata. Pruning occurs within a stratum (for which observations are indistinguishable to `cem` proper) by random matching inside `cem` strata⁵.

Here is an example of this approach. Running the earlier call with the `k2k` options yields:

```
. cem age education black nodegree re74, tr(treated) k2k
```

Matching Summary:

Number of strata: 205

Number of matched strata: 67

	0	1
All	425	297
Matched	205	205
Unmatched	220	92

Multivariate L1 distance: .37560976

Univariate imbalance:

	L1	mean	min	25%	50%	75%	max
age	.07805	-.10732	0	0	0	0	-1
education	0	0	1	0	0	0	0
black	0	0	0	0	0	0	0
nodegree	0	0	0	0	0	0	0
re74	.0439	-34.547	0	0	-120.7	-214.55	96.881

It is clear that the number of matched units has decreased after using the `k2k` option.

4.4 Estimating the Causal Effect from `cem` output

Using the output from `cem`, we can estimate the SATT by the regular Stata methods, by simply including the `cem_weights`. For example,

```
. reg re78 treated [iweight=cem_weights]
```

Source	SS	df	MS	Number of obs =	552
Model	128314324	1	128314324	F(1, 550) =	3.15
Residual	2.2420e+10	550	40764521.6	Prob > F =	0.0766
Total	2.2549e+10	551	40923414.2	R-squared =	0.0057
				Adj R-squared =	0.0039
				Root MSE =	6384.7

re78	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
------	-------	-----------	---	------	----------------------

⁵Note that in the R version of this software pruning within strata can be done using a distance metric.

treated	979.1905	551.9132	1.77	0.077	-104.9252	2063.306
_cons	4919.49	354.7061	13.87	0.000	4222.745	5616.234

For convenience, we compute this as a regression of the outcome variable on a constant and the treatment variable, where the SATT estimate is the coefficient on the `treated` variable, in our case 979.19. Any Stata command that accepts weights (`aweight` or `iweight`) can be used.

If exact matching (i.e., without coarsening) was chosen this procedure is appropriate as is. In other situations, with some coarsening, some imbalance remains in the matched data. The remaining imbalance is strictly bounded by the level of coarsening, which can be seen by any remaining variation within the coarsened bins. Thus, a reasonable approach in this common situation is to attempt to adjust for the remaining imbalance via a statistical model. (Modeling assumptions for models applied to the matched data are much less consequential than they would otherwise be because CEM is known to strictly bound the level of model dependence.) To apply a statistical model to control for the remaining imbalance, we simply add variables to the regression command. For example:

```
. reg re78 treated re74 re75 [iweight=cem_weights]
```

Source	SS	df	MS	Number of obs =	552
Model	649651702	3	216550567	F(3, 548) =	5.42
Residual	2.1899e+10	548	39961951.7	Prob > F =	0.0011
Total	2.2549e+10	551	40923414.2	R-squared =	0.0288
				Adj R-squared =	0.0235
				Root MSE =	6321.5

re78	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
treated	988.083	546.5395	1.81	0.071	-85.48584 2061.652
re74	-.0174322	.1593346	-0.11	0.913	-.3304134 .2955491
re75	.3190651	.1744905	1.83	0.068	-.023687 .6618172
_cons	4287.523	393.0883	10.91	0.000	3515.378 5059.667

The user can also specify `glm` modeling in the case of binary, count, or other non-continuous outcome variables by utilizing their commands in Stata (`logit`, `poisson`, etc) combined with the `iweight` syntax.

4.5 Matching and Missing Data

Almost all previous methods of matching assume the absence of any missing values. In contrast, CEM offers two approaches to dealing with missing values (item nonresponse). In the first, where we treat missing values as one of the values of the variables, is appropriate when “.” is a valid value that is not really missing (such as when “no opinion” really means no opinion). The other is a special procedure to allow for multiply imputed data in CEM.

4.5.1 Matching on Missingness

If users leave missing values in the data, `cem` will coarsen the variables as normal, but use “.” as a separate category for each variable. Thus, `cem` will match on missingness.

4.5.2 Matching Multiply Imputed Data in a Single File

Consider a data set to be matched, some of which is missing. One approach to analyzing data with missing values is *multiple imputation*, which involves creating m (usually about $m = 5$) data sets, each of which is the same as the original except that the missing values have been imputed in each. Uncertainty in the values of the missing cells is represented by variation in the imputations across the different imputed data sets (King et al., 2001).

When one uses the `mi` commands in Stata to impute missing data, the program can return the data in a stacked or `flong` format. The `cem` command can take advantage of this format to easily match multiply imputed datasets. Note that the data must be in the `flong` format: one dataset with $\tilde{n} = n + m \times n$ observations, where m is the number of imputed datasets and n is the number of observations in the original dataset. The first n rows of the data are the original dataset and the following $m \times n$ rows are each of the imputed datasets. When using the `mi` commands in Stata, you can ensure the data are in this format by using `mi set flong` or `mi convert flong`.

Once you have your imputed data in `flong` format, you can pass the data to `cem` along with the imputation variable and `cem` will match the multiple imputations. The CEM algorithm first runs on the entire stacked dataset of imputations, with possibly different strata for each imputation. To combine strata across imputation, CEM chooses the strata most often assigned to an observation. This strata assignment is given to each of the imputed datasets (that is, the `cem_weights` variable is added to each of the datasets).

To run this command simply add the `impvar` option:

```
. use "http://www.mattblackwell.org/files/stata/imp-lelonde.dta", clear  
(Written by R. )
```

```
. mi import flong, m(imp) id(mi_id)
```

```
. cem age education black nodegree re74, tr(treated) impvar(imp)  
(using the scott break method for imbalance)
```

Matching Summary:

Number of strata: 244

Number of matched strata: 85

	0	1
All	425	297
Matched	299	229
Unmatched	126	68

Multivariate L1 distance: .37560605

Univariate imbalance:

	L1	mean	min	25%	50%	75%	max
age	.08093	-.05183	0	0	0	0	-1
education	.01276	-.00457	0	0	0	0	0
black	.00218	.00063	0	0	0	0	0
nodegree	.00437	-.00092	0	0	0	0	0
re74	.05257	1.1624	0	0	104.2	232.05	1401.7

Now that we have run `cem`, we can then use `mi estimate` to easily combine regressions from the imputed datasets to estimate the SATT (King et al., 2001), being sure to weight by the `cem_weights`:

```
. mi estimate: reg re78 treated [iweight=cem_weights]
```

Multiple-imputation estimates	Imputations	=	5
Linear regression	Number of obs	=	528
	Average RVI	=	0.0353
	Largest FMI	=	0.0308
	Complete DF	=	526
DF adjustment: Small sample	DF: min	=	455.20
	avg	=	472.43
	max	=	489.65
Model F test: Equal FMI	F(1, 455.2)	=	4.04
Within VCE type: OLS	Prob > F	=	0.0450

re78	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
treated	1112.818	553.5378	2.01	0.045	25.01139	2200.624
_cons	4867.038	362.5507	13.42	0.000	4154.691	5579.385

In addition to the built-in multiple imputation tools, there are several user-written tools for combining multiple regression output such as `miest` written by Ken Scheve⁶ and `clarify`⁷.

4.5.3 Matching Multiply Imputed Data in Separate Files

Suppose that we have used some imputation program (such as (Honaker, King and Blackwell, 2010)) to produce 5 imputed datasets, saved as:

```
imp1.dta
imp2.dta
imp3.dta
imp4.dta
imp5.dta
```

⁶`miest` is available at <http://gking.harvard.edu/amelia/amelia1/docs/mi.zip>.

⁷`Clarify` is available at <http://gking.harvard.edu/stats.shtml#clarify>

As an example we added missingness to the example dataset and imputed it using Honaker, King and Blackwell (2010)⁸. If we place all of the imputed datasets in the same directory and open the first, we can run `cem` with the `miname` and `misets` arguments to specify the root of the imputed datasets' filename and the number of datasets, respectively. In our example, this would be:

```
. use imp1.dta, clear

.
. cem age education black nodegree re74, tr(treated) miname(imp) misets(5)
```

Matching Summary:

```
-----
Number of strata: 235
Number of matched strata: 76
```

	0	1
All	425	297
Matched	312	217
Unmatched	113	80

Multivariate L1 distance: .38286064

Univariate imbalance:

	L1	mean	min	25%	50%	75%	max
age	.02132	-.07344	.19196	0	1	0	-1
education	.01173	-.0121	1	0	0	0	0
black	.00207	.00041	0	0	0	0	0
nodegree	.00461	-.00092	0	0	0	0	0
re74	.04987	-4.1404	-398.68	0	375.1	-236.7	96.881

Note that `cem` outputs the data in a single, stacked format so that we can use the estimation procedure as with a single file above.

4.6 Blocking in Randomized Experiments

CEM can produce strata for a block randomized design for a set of pre-treatment covariates. As block randomized designs outperform complete randomization on bias, efficiency, power and robustness, it should be used whenever possible (Imai, King and Nall, 2009; Imai, King and Stuart, 2008). To create a set of strata for a block randomized design, simply run CEM without passing a treatment variable. This will assign observations to strata based on their coarsened values and create a `cem_strata` variable, indicating this assignment. Once this is complete, simply randomly assign treatment within these strata to complete the block randomized design.

⁸If users are interested in working with this example, they can access these sample files at <http://gking.harvard.edu/cem/imp1.dta>, etc. Once all five are downloaded, users can generate the following output. The original data file with missingness added is at <http://gking.harvard.edu/cem/lelonde.dta>

4.7 Using cem to Improve Other Matching Methods

Even if you plan to use a different matching method, you can still use the CEM algorithm to improve that matching solution. An important step before matching is restricting the data to areas of common empirical support. This avoids making inference based on extrapolation as such inferences are known to be extremely model dependent. Traditional matching methods, however, are not equipped to handle this situation. For example, the propensity score can be used to find the area of extrapolation only after we know that the correct propensity score model has been used. However, the only way to verify that the correct propensity score model has been specified is to check whether matching on it produces balance between the treated and control groups on the relevant covariates. But balance cannot be reliably checked until the region of extrapolation has been removed. To avoid this type of infinite regress, researchers use entirely different technologies for the first step, such as kernel density estimation (Heckman, Ichimura and Todd, 1998) or dropping control units outside the hyper-rectangle (Iacus and Porro, 2009) or convex hull (King and Zeng, 2006) of the treated units.

The matching methods currently in Stata all rely on propensity score methods for restricting the data to common empirical support. For CEM, on the other hand, this restriction is a natural consequence of the algorithm. All observations within a stratum containing both a treated and control unit are by definition inside of the common support. In light of this, a good use of CEM would be to reduce the data to common support before applying another matching solution such as `psmatch2`, `nnmatch`, or `pscore`. This will improve the quality of the inferences drawn from these methods. Once you have run `cem`, all you must do is run the following command to restrict the data to common support:

```
. drop if cem_matched == 0
```

Alternatively, you can use any of the matching methods with an `if cem_matched==1` option. This will force the other matching methods to only match in the region of common support. As an example using `nnmatch`, this would be

```
. nnmatch re78 treated age education black nodegree re74 if cem_matched == 1
```

Of course you can apply this idea to any matching method in Stata, not just the ones listed here.

5 cem — Coarsened Exact Matching

5.1 Syntax

```
. cem varname1 [(cutpoints1)] [varname2 [(cutpoints2)]] ... [if] [in] [, treatment(varname)
showbreaks autocuts(string) k2k imbbreaks(string) miname(string) misets(#)]
```

5.2 Description

`cem` implements the Coarsened Exact Matching method described in Iacus, King and Porro (2012). The main inputs for `cem` are the variables to use (*varname#*) and the cutpoints that define the coarsening (*cutpoints#*). The latter option is set in a parentheses after the name of the relevant variable. Users can either specify cutpoints for a variable or allow `cem` to automatically coarsen the data based on an automatic coarsening algorithm, chosen by the user. To specify a set of cutpoints for a variable, place a numlist in parentheses

after the variable's name. To specify an automatic coarsening, place a string indicating the binning algorithm to use in parentheses after the variable's name. To create a certain number of equally spaced cutpoints including the extreme values, say 10, place **#10** in the parentheses (using **#0** will force **cem** into not coarsening the variable at all). Omitting the parenthetical statement after the variable name tells **cem** to use the default binning algorithm, itself set by **autocuts**.

Note that character variables are ignored by **cem**. These variables will need to be converted into numeric variables using **encode**. Coarsening that are not ordinal must be done before running **cem** using the **recode** command, as described above.

5.3 Options

treatment(varname) tells **cem** which variable should be used for matching.

showbreaks will have **cem** display the cutpoints used for each variable on the screen.

autocuts(string) sets the default automatic coarsening algorithm. The default for this is "sturges". Any variable without a **cutpoint#** command after its name will use the **autocuts** argument.

k2k will have **cem** produce a matching result that has the same number of treated and control in each matched strata by randomly dropping observations.

imbbreaks(string) sets the coarsening method for the imbalance checks printed after **cem** runs. This should match whichever method is used for imbalance checks elsewhere. If either **cem** or **imbhas** been run and there is a **r(L1.breaks)** available, this will be the default. Otherwise, the default for this is "scott"

miname(string) is the root of the filenames of the imputed dataset. They should be in the working directory. For example, if **miname** were "imputed", then the filenames should be "imputed1.dta", "imputed2.dta" and so on.

misets(#) is the number of imputed datasets being used for matching.

5.4 Output

The output **cem** returns depends on the inclusion of a treatment variable. If the treatment variable is provided, **cem** will match and return the following three variables in the current dataset:

cem_strata is the strata number assigned to each observation by **cem**.

cem_matched is 1 for a matched observation and 0 for an unmatched observation.

cem_weights is the weight of the stratum for each observation. Strata with unmatched units are give 0 weight and treated observations are given a weight of 1.

cem_treat when using the multiple imputation features, **cem** outputs this variable, which is the treatment vector used for matching. **cem** applies the same combination rule to treatment as to strata.

If the options for multiple imputation are used, **cem** saves each of these variables in each of the imputed datasets, allowing for easy use in programs like **miest**.

The following are stored as saved results in Stata's memory:

Scalars

<code>r(n_strata)</code>	number of strata.
<code>r(n_groups)</code>	number of treatment levels.
<code>r(n_mstrata)</code>	number of strata with matches.
<code>r(n_matched)</code>	number of matched observations.
<code>r(L1)</code>	multivariate imbalance measure.

Matrices

<code>r(match_table)</code>	table of treatment vs matched.
<code>r(groups)</code>	tabulation of treatment variable.
<code>r(imbal)</code>	univariate imbalance measures.

Macros

<code>r(varlist)</code>	covariate variables used.
<code>r(treatment)</code>	treatment variable.
<code>r(cem_call)</code>	call to cem.
<code>r(L1_breaks)</code>	break method used for L1 distance

If the treatment variable is omitted (e.g. for blocking), then the only outputs are `cem_strata`, `r(n_strata)`, `r(varlist)`, and `r(cem_call)`.

6 imb - Imbalance Measures for CEM

6.1 Syntax

`imb varlist [if] [in] [, treatment(varname) breaks(string) miname(string) misets(string) useweights]`

6.2 Description

`imb` returns a number of measures of imbalance in covariates between treatment and control groups. A multivariate L1 distance, univariate L1 distances, difference in means and empirical quantiles difference are reported. The L1 measures are computed by coarsening the data according to `breaks` and comparing across the multivariate histogram. See Iacus, King and Porro (2012) for more details on this measure.

6.3 Arguments

`treatment(varname)` sets the treatment variable used for the imbalance checks.

`breaks(string)` sets the default automatic coarsening algorithm. If either `cem` or `imb` has been run and there is a `r(L1_breaks)` available, this will be the default. Otherwise, the default for this is “scott”. It is not incredibly important which method is used here as long as it is consistent.

`miname(string)` is the root of the filenames of the imputed dataset. They should be in the working directory. For example, if `miname` were “imputed”, then the filenames should be “imputed1.dta”, “imputed2.dta” and so on.

`misets(integer)` is the number of imputed datasets being used for matching.

`useweights` makes `imb` use the weights from the output of `cem`. This is useful for checking balance after running `cem`.

6.4 Saved Results

Scalars

`r(L1)` multivariate imbalance measure

Matrices

`r(imbal)` matrix of univariate imbalance measures

Macros

`r(L1_breaks)` break method used for L1 distance

References

- Abadie, Alberto, Alexis Diamond and Jens Hainmueller. 2009, forthcoming. “Synthetic Control Methods for Comparative Case Studies: Estimating the Effect of California’s Tobacco Control Program.”
- Abadie, Alberto, David Drukker, Jane Leber Herr and Guido W. Imbens. 2004. “Implementing Matching Estimators for Average Treatment Effects in Stata.” *The Stata Journal* 4(3):290–311.
- Becker, Sascha O. and Andrea Ichino. 2002. “Estimation of average treatment effects based on propensity scores.” *The Stata Journal* 2(4):358–377.
- Crump, Richard K., V. Joseph Hotz, Guido W. Imbens and Oscar Mitnik. 2006. “Moving the Goalposts: Addressing Limited Overlap in Estimation of Average Treatment Effects by Changing the Estimand.” Department of Economics, UC Berkeley.
- Heckman, James, H. Ichimura and P. Todd. 1998. “Matching as an Econometric Evaluation Estimator: Evidence from Evaluating a Job Training Program.” *Review of Economic Studies* 65:261–294.
- Ho, Daniel, Kosuke Imai, Gary King and Elizabeth Stuart. 2007. “Matching as Nonparametric Preprocessing for Reducing Model Dependence in Parametric Causal Inference.” *Political Analysis* 15:199–236. <http://gking.harvard.edu/files/abs/matchp-abs.shtml>.
- Honaker, James, Gary King and Matthew Blackwell. 2010. “Amelia II: A Program for Missing Data.” <http://gking.harvard.edu/amelia>.
- Iacus, Stefano M., Gary King and Giuseppe Porro. 2012. “Causal Inference Without Balance Checking: Coarsened Exact Matching.” *Political Analysis* . <http://gking.harvard.edu/files/abs/cem-plus-abs.shtml>.
- Iacus, Stefano M. and Giuseppe Porro. 2009. “Random Recursive Partitioning: a matching method for the estimation of the average treatment effect.” *Journal of Applied Econometrics* 24:163–185.
- Imai, Kosuke, Gary King and Clayton Nall. 2009. “The Essential Role of Pair Matching in Cluster-Randomized Experiments, with Application to the Mexican Universal Health Insurance Evaluation.” *Statistical Science* 24(1):29–53. <http://gking.harvard.edu/files/abs/cluster-abs.shtml>.
- Imai, Kosuke, Gary King and Elizabeth Stuart. 2008. “Misunderstandings Among Experimentalists and Observationalists about Causal Inference.” *Journal of the Royal Statistical Society, Series A* 171, part 2:481–502. <http://gking.harvard.edu/files/abs/matchse-abs.shtml>.
- King, Gary, James Honaker, Anne Joseph and Kenneth Scheve. 2001. “Analyzing Incomplete Political Science Data: An Alternative Algorithm for Multiple Imputation.” *American Political Science Review* 95(1, March):49–69. <http://gking.harvard.edu/files/abs/evil-abs.shtml>.
- King, Gary and Langche Zeng. 2006. “The Dangers of Extreme Counterfactuals.” *Political Analysis* 14(2):131–159. <http://gking.harvard.edu/files/abs/counterft-abs.shtml>.

- Lalonde, Robert. 1986. "Evaluating the Econometric Evaluations of Training Programs." *American Economic Review* 76:604–620.
- Leuven, Edwin and Barbara Sianesi. 2004. "PSMATCH2: Stata module to perform full Mahalanobis and propensity score matching, common support graphing, and covariate imbalance testing." EconPapers. <http://econpapers.repec.org/software/bocbocode/S432001.htm>.